

13. Loading and Preprocessing Data with TensorFlow

What we'll cover:

- The Data API
- The TFRecord Format
- Preprocessing the Input Features
- TF Transform
- The TensorFlow Datasets (TFDS) Project

The Data API

Create dataset object

Tell it where to get the data (text, binary, SQL, etc.)

...and how to transform it

Processing - write your own, or use standard layers in Keras:

Normalization

Encoding of non-numerical features

The Data API

Dataset

- Usually reads gradually from disk, but can operate entirely in RAM

```
tf.data.Dataset.from_tensor_slices(<any data tensor>)
```

- Slices of <tensor> along the 1st dimension

The Data API - Chaining Transformations

```
X = tf.range(10) # any data tensor
dataset = tf.data.Dataset.from_tensor_slices(X)
dataset = dataset.repeat(3).batch(7)
for item in dataset:
    . . . <do something>
```


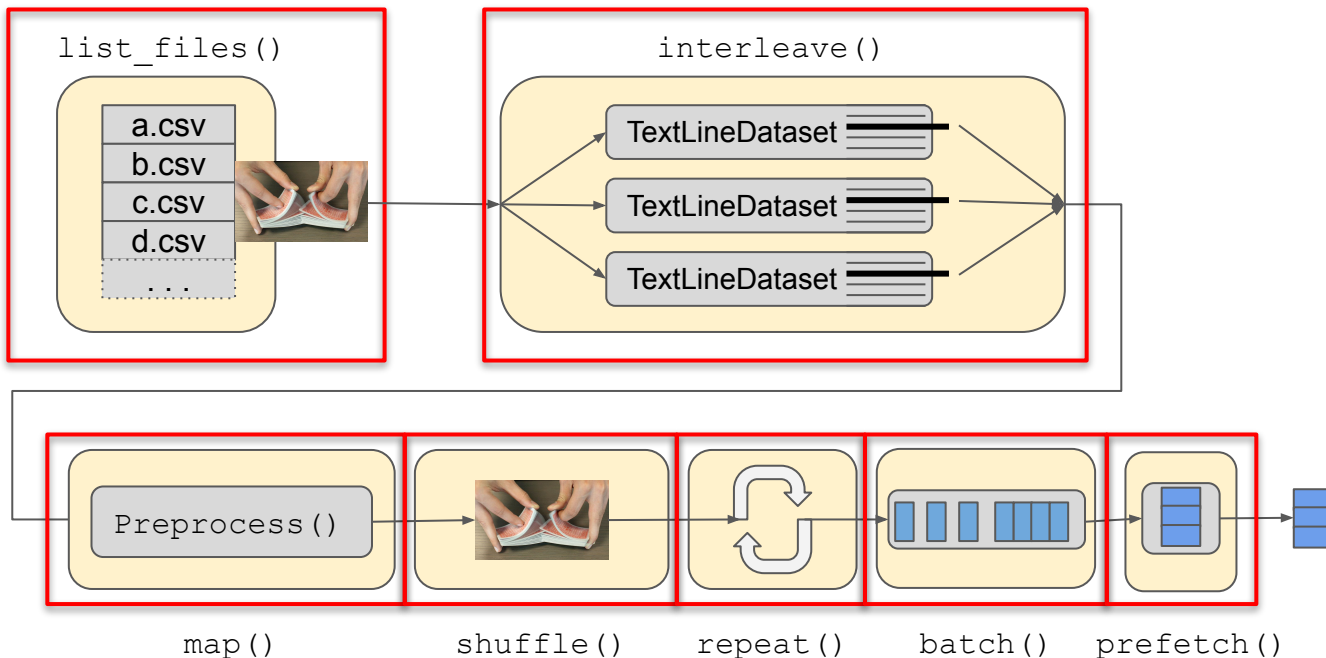


Figure 13-1. Chaining dataset transformations

The Data API - other functions

```
# must be convertible to a TF function
# apply to each item
dataset = dataset.map(lambda x: x * 2)
# apply to entire dataset
dataset = dataset.apply(tf.data.experimental.unbatch())
```



```
def csv_reader_dataset (filepaths, repeat=1, n_readers=5, n_read_threads=None,
shuffle_buffer_size=10000, n_parse_threads=5, batch_size=32):
    dataset = tf.data.Dataset.list_files(filepaths)
    dataset = dataset.interleave(lambda filepath: tf.data.TextLineDataset(filepath).skip(1), \
                                cycle_length=n_readers, num_parallel_calls=n_read_threads)
    dataset = dataset.map(preprocess, num_parallel_calls=n_parse_threads)
    dataset = dataset.shuffle(shuffle_buffer_size).repeat(repeat)

    return dataset.batch(batch_size).prefetch(1)
```

Preprocessing the Input Features

```
X_mean, X_std = [ ... ] # mean and scale of each feature in the training set

n_inputs = 8

def preprocess(line):
    defs = [0.] * n_inputs + [tf.constant([], dtype=tf.float32)]
    fields = tf.io.decode_csv(line, record_defaults=defs)
    x = tf.stack(fields[:-1])
    y = tf.stack(fields[-1:])
    return (x - X_mean) / X_std, y
```


Preprocessing the Input Features

```
class Standardization(keras.layers.Layer):  
    def adapt(self, data_sample):  
        self.means_ = np.mean(data_sample, axis=0, keepdims=True)  
        self.stds_ = np.std(data_sample, axis=0, keepdims=True)  
    def call(self, inputs):  
        return(inputs - self.means_) / (self.stds_ + keras.backend.epsilon())
```

```
std_layer = Standardization()  
std_layer.adapt(data_sample)
```

```
model = keras.Sequential()  
model.add(std_layer) [ ... ] # create the rest of the model  
model.compile([ ... ])  
model.fit([ ... ])
```

```
# Or  
keras.layers.Normalization([...])
```

```
# see https://keras.io/api/layers/normalization\_layers/
```

The TFRecord Format

```
with tf.io.TFRecordWriter("my_data.tfrecord") as f:  
    f.write(b"This is the first record")  
    f.write(b"And this is the second record")
```

Protobufs

```
syntax = "proto3";  
message ByteList { repeated bytes value = 1; }  
message FloatList { repeated float value = 1 [packed = true]; }  
message Int64List { repeated int64 value = 1 [packed = true]; }  
message Feature {  
    oneof kind {  
        ByteList bytes_list = 1;  
        FloatList float_list = 2;  
        Int64List int64_list = 3;  
    }  
};  
message Features { map<string, Feature>feature = 1; };  
message Example { Features features = 1; };
```

The TFRecord Format - Example Protobuf

```
from tensorflow.train import ByteList, FloatList, Int64List
from tensorflow.train import Feature, Features, Example

person_example = Example(
    features = Features(feature = {
        "name": Feature(bytes_list=ByteList(value=[b"Alice"])),
        "id": Feature(int64_list=Int64List(value=[123])),
        "emails": Feature(bytes_list=ByteList(value=[b"a@b.com", b" c@d.com"]))
    })

# Writing
with tf.io.TFRecordWriter("my_contacts.tfrecord") as f:
    f.write(person_example.SerializeToString())
```

The TFRecord Format - reading

```
# Reading
feature_description = {
    "name": tf.io.FixedLenFeature([], tf.string, default_value=""),
    "id": tf.io.FixedLenFeature([], tf.int64, default_value=0),
    "emails":tf.io.VarLenFeature(tf.string),
}

for serialized_example in tf.data.TFRecordDataset(["my_contacts.tfrecord"]):
    parsed_example = tf.io.parse_single_example(serialized_example, feature_description)
```

The TFRecord Format - SequenceExample Protobuf

```
message FeatureList { repeated Feature feature = 1; };  
message FeatureLists { map<string, FeatureList>feature_list = 1; };  
message SequenceExample {  
    Features context = 1;  
    FeatureLists feature_lists = 2;  
};
```

```
# Reading  
parsed_context , parsed_feature_lists = tf.io.parse_single_sequence_example( \  
    serialized_sequence_example, context_feature_descriptions, sequence_feature_descriptions)  
parsed_content = tf.RaggedTensor.from_sparse(parsed_feature_lists["content"])
```

TF Transform (tf.Transform)

```
import tensorflow_transform as tft

def preprocess(inputs): # inputs = a batch of input features
    median_age = inputs["housing_median_age"]
    ocean_proximity = inputs["ocean_proximity"]
    standardized_age = tft.scale_to_z_score(median_age)
    ocean_proximity_id = tft.compute_and_apply_vocabulary(ocean_proximity)
    return {
        "standardized_median_age": standardized_age,
        "ocean_proximity_id": ocean_proximity_id
    }
```

The TensorFlow Datasets (TFDS) Project

```
import tensorflow_datasets as tfds

dataset = tfds.load(name = "mnist", batch_size=32, as_supervised=True)
mnist_train = dataset["train"].prefetch(1)
model = keras.models.Sequential([ ... ])
model.compile(loss="sparse_categorical_crossentropy", optimizer="sgd")
model.fit(mnist_train, epochs=5)
```

Questions?

THANK YOU!