

# Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow, 2<sup>nd</sup> edition by Aurélien Geron Chapter 11

San Diego Machine Learning  
2022 JAN 15  
Discussion Leader: Robert Kraig

# CHAPTER 11 OUTLINE

- ▶ Vanishing / Exploding Gradients
- ▶ Reusing Pretrained Layers
- ▶ Faster Optimizers
- ▶ Regularization to Avoid Overfitting

```
tf.keras.layers.Dense(  
    units, activation=None, use_bias=True,  
    kernel_initializer='glorot_uniform',  
    bias_initializer='zeros', kernel_regularizer=None,  
    bias_regularizer=None, activity_regularizer=None, kernel_constraint=None,  
    bias_constraint=None, **kwargs  
)
```

```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),  
              loss=tf.keras.losses.BinaryCrossentropy(),  
              metrics=[tf.keras.metrics.BinaryAccuracy(),  
                      tf.keras.metrics.FalseNegatives()])
```

TensorFlow  
Hub

## 5.10 Building a Machine Learning Algorithm

Nearly all deep learning algorithms can be described as particular instances of a fairly simple recipe: combine a specification of a dataset, a cost function, an optimization procedure and a model.

*Deep Learning*, Goodfellow et al.

# 1. Vanishing and Exploding Gradients

- Glorot and Bengio (2010):
  - Vanishing Gradients Expected in early NNs
    - Standard Normal Initialization
    - Sigmoid Activation: Small Derivatives


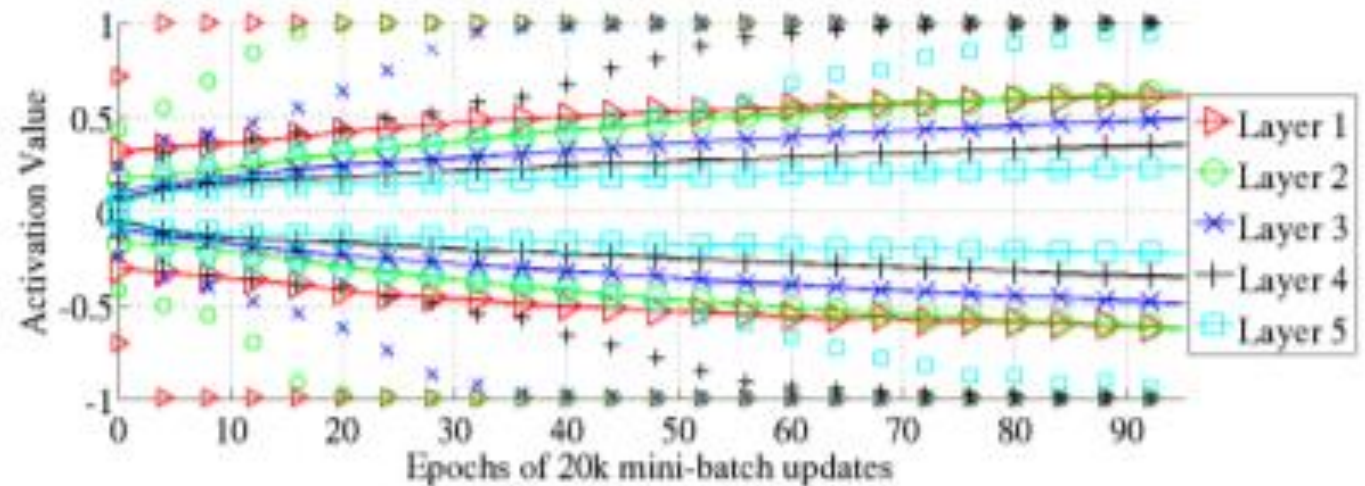
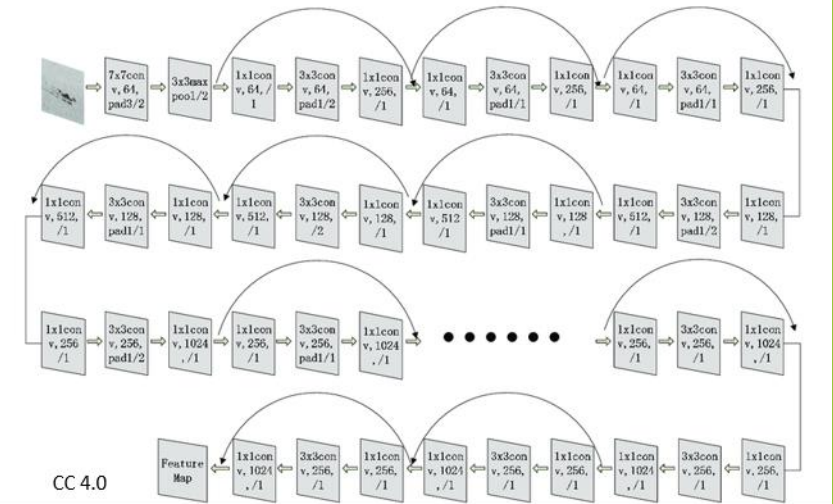


Figure 11-1. Logistic activation function saturation



Glorot and Bengio

# Initializers

- ▶ Glorot and Bengio observed that, for stability:
  - ▶ Each layer should have equal variance of input and output activations (Forward)
  - ▶ Each layer should have equal variance of input and output gradients (Back-Prop)

## Equation 11-1. Glorot initialization (when using the logistic activation function)

Normal distribution with mean 0 and variance  $\sigma^2 = \frac{1}{fan_{avg}}$

Or a uniform distribution between  $-r$  and  $+r$ , with  $r = \sqrt{\frac{3}{fan_{avg}}}$

$$fan_{avg} = (fan_{in} + fan_{out})/2$$

Table 11-1. Initialization parameters for each type of activation function

Initialization	Activation functions	$\sigma^2$ (Normal)
Glorot	None, tanh, logistic, softmax	$1 / fan_{avg}$
He	ReLU and variants	$2 / fan_{in}$
LeCun	SELU	$1 / fan_{in}$

# Non-Saturating Activation Functions





Figure 11-2: Leaky ReLU: like ReLU, but with a small slope for negative values





Figure 11-3: ELU activation function



# Batch Normalization

- ▶ [Ioffe and Szegedy, 2015](#)
- ▶ Prevent Gradient Problems during Training
- ▶ Each Neuron (or Layer if CNN)
- ▶ Why does it work?
  - ▶ Internal Covariate Shift, Gradient Smoothing
  - ▶ Regularization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$


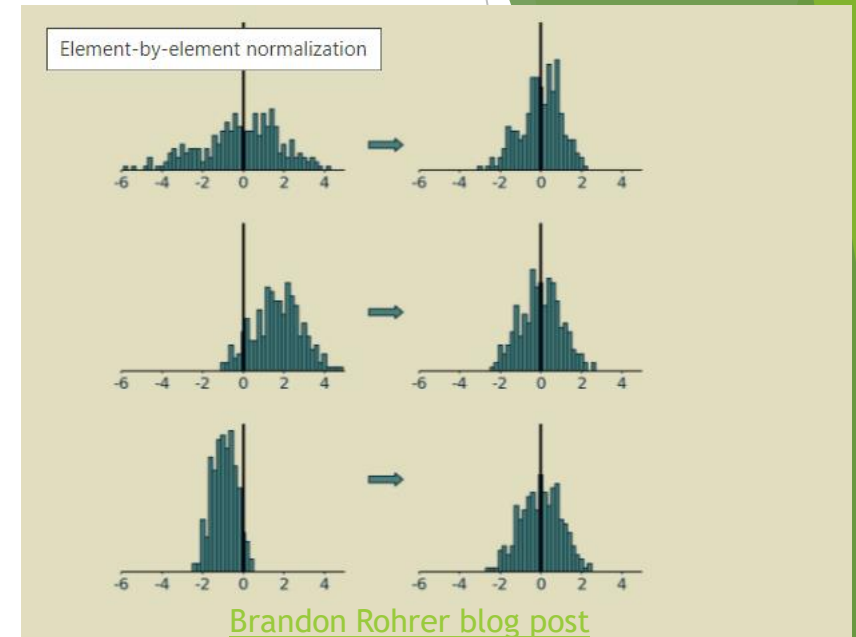
$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.



# Implementing Batch Norm w/ Keras

- ▶ Where to put BatchNorm
  - ▶ Typical: Separate Layer after Activation
  - ▶ Can do it before Activation instead (a la Ioffe and Szegedy)
- ▶ Parameters: Trainable vs Non-Trainable
- ▶ Inference vs Training
  - ▶ Training: separate mean/std calculation on each batch
  - ▶ Inference: Estimate dataset mean/std using momentum

$$\hat{\mathbf{v}} \leftarrow \hat{\mathbf{v}} \times \text{momentum} + \mathbf{v} \times (1 - \text{momentum})$$

- ▶ BN is Ubiquitous now
  - ▶ Omitted in diagrams / papers
  - ▶ But is it necessary? ... (Fixed update initialization) <https://arxiv.org/pdf/1901.09321.pdf>

# Gradient Clipping

- ▶ Exploding gradients don't matter if you clip them during BackProp

```
optimizer = keras.optimizers.SGD(clipvalue=1.0)  
model.compile(loss="mse", optimizer=optimizer)
```

- ▶ Used especially in RNNs
- ▶ Use clipnorm instead of clipvalue to preserve direction (L2 norm)



## 2. Reusing Pretrained Layers

- ▶ aka => Transfer Learning
- ▶ Freeze Lower Layers
- ▶ Retrain Upper Layers
- ▶ Replace Input/Output as Needed

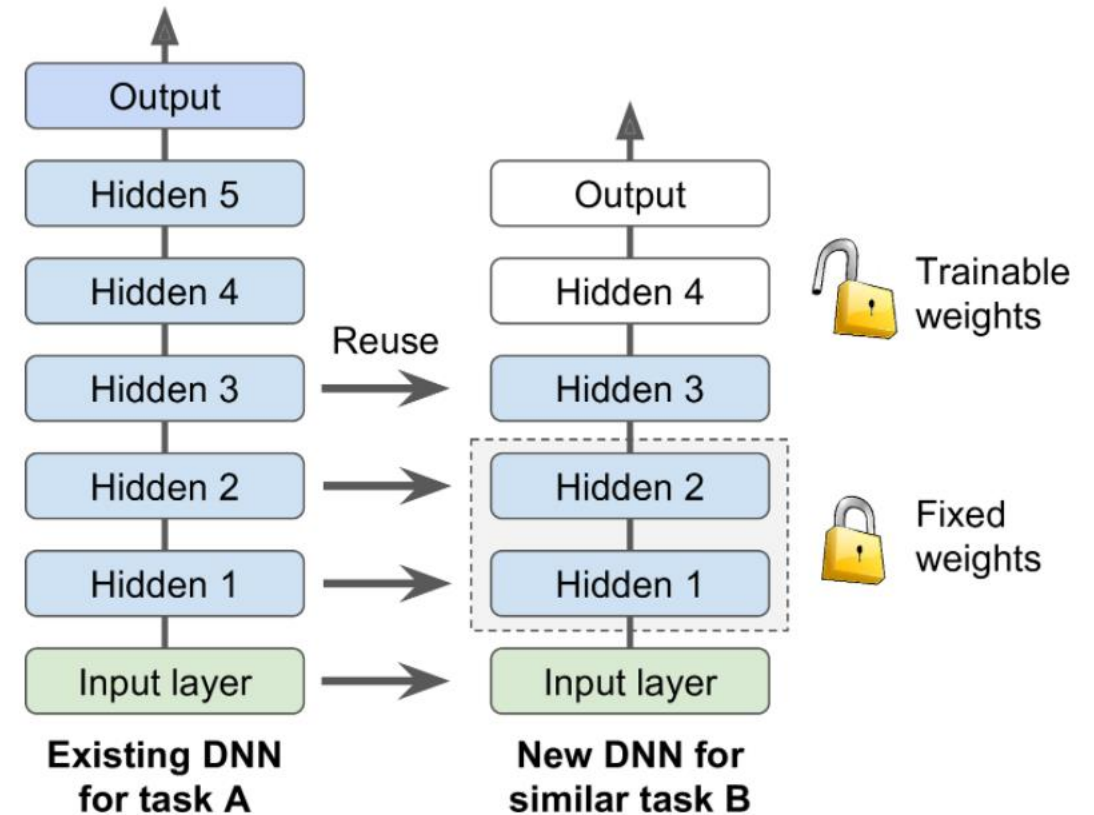



Figure 11-4. Reusing pretrained layers

# Transfer Learning with Keras

## ► Fashion MNIST

```
model_A = keras.models.load_model("my_model_A.h5")  
model_B_on_A = keras.models.Sequential(model_A.layers[:-1])  
model_B_on_A.add(keras.layers.Dense(1, activation="sigmoid"))
```

- Clone avoids model\_A overwrite
- One approach: train in stages.
  - Freeze middle layers at first
  - Train new/upper layers
  - Tweak middle after stabilization
- Geron Fashion MNIST example: Confession under Torture
- Transfer Learning works best with large/sparse models



# Unsupervised Pretraining

- Data is cheap. Labels are expensive!
- Option: use Autoencoding
- Hinton (2006): Greedy Layer-Wise Pretraining




Figure 17-1. The chess memory experiment (left) and a simple autoencoder (right)

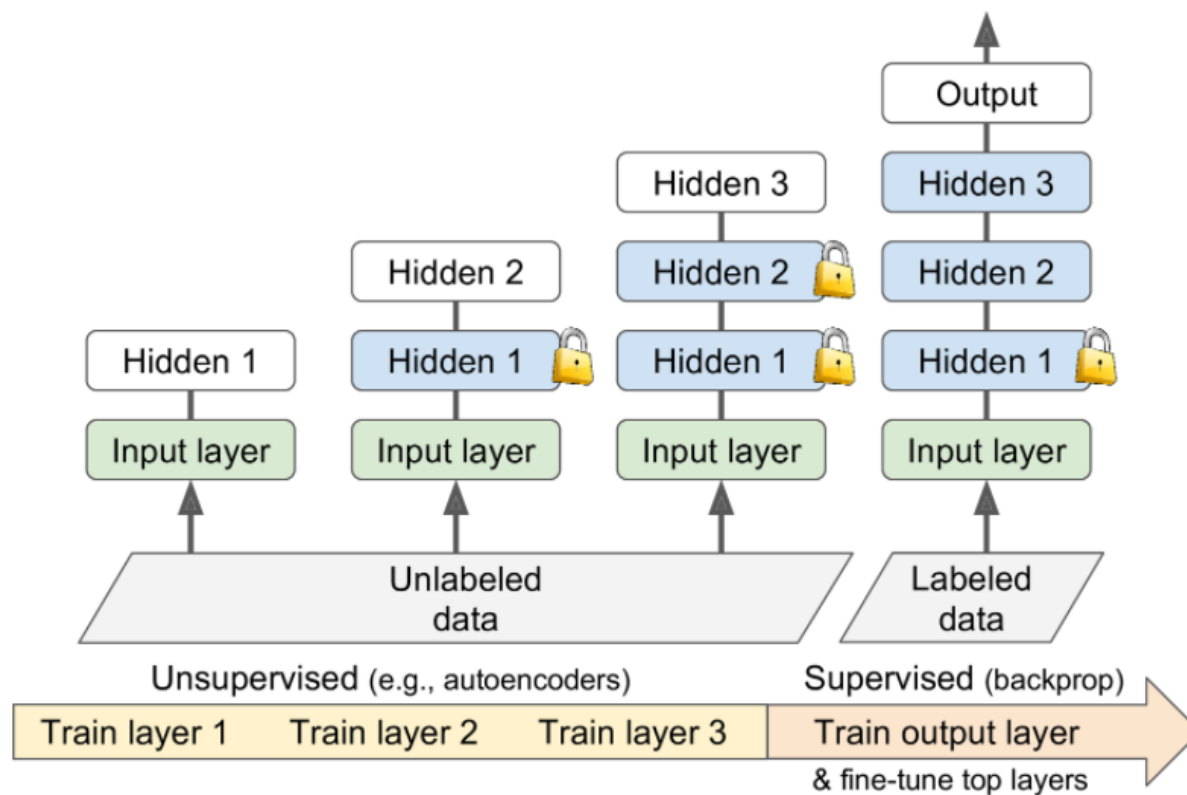


Figure 11-5. In unsupervised training, a model is trained on the unlabeled data (or on all the data) using an unsupervised learning technique, then it is fine-tuned for the final task on the labeled data using a supervised learning technique; the unsupervised part may train one layer at a time as shown here, or it may train the full model directly


# Pretraining on an Auxiliary Task

- ▶ Computer Vision Example
  - ▶ Face Classifier from very few images
  - ▶ Pull face data of random people from web
  - ▶ Train binary classifier for image pairs: is it the same person?
- ▶ NLP
  - ▶ Language models (e.g. BERT) are trained with masking
  - ▶ Can reuse latent representations for many NLP tasks

# Faster Optimizers: Momentum

- Polyak 1964: [Some methods of speeding up the convergence of iteration methods](#)
- Gabriel Goh (2017): [Interactive Visualization of Momentum](#)

## Why Momentum Really Works



Step-size  $\alpha = 0.02$



Momentum  $\beta = 0.99$



We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?

Equation 11-4. Momentum algorithm

1.  $\mathbf{m} \leftarrow \beta \mathbf{m} - \eta \nabla_{\theta} J(\theta)$
2.  $\theta \leftarrow \theta + \mathbf{m}$

```
optimizer = keras.optimizers.SGD(lr=0.001, momentum=0.9)
```

# Faster Optimizers: Nesterov Accelerated Gradient

## ► Yurii Nesterov (1983)

Equation 11-5. Nesterov Accelerated Gradient algorithm

1.  $\mathbf{m} \leftarrow \beta \mathbf{m} - \eta \nabla_{\theta} J(\theta + \beta \mathbf{m})$
2.  $\theta \leftarrow \theta + \mathbf{m}$




Figure 11-6. Regular versus Nesterov momentum optimization: the former applies the gradients computed before the momentum step, while the latter applies the gradients computed after

```
optimizer = keras.optimizers.SGD(lr=0.001, momentum=0.9, nesterov=True)
```

# Faster Optimizers: AdaGrad

- ▶ Duchi et al: Adaptive Subgradient Methods for Online Learning and Stochastic Optimization (2011)
- ▶ Adaptive Learning Rate
  - ▶ Different across dimensions
- ▶ Good for quadratics
- ▶ Stops early in complex problems

Equation 11-6. AdaGrad algorithm

1.  $\mathbf{s} \leftarrow \mathbf{s} + \nabla_{\theta} J(\theta) \otimes \nabla_{\theta} J(\theta)$
2.  $\theta \leftarrow \theta - \eta \nabla_{\theta} J(\theta) \oslash \sqrt{\mathbf{s} + \varepsilon}$

```
optimizer = keras.optimizers.Adagrad(learning_rate=0.001)
```




Figure 11-7. AdaGrad versus Gradient Descent: the former can correct its direction earlier to point to the optimum

# Faster Optimizers: RMSProp

- ▶ RMSProp improves AdaGrad
- ▶ Adds Exponential Decay to the first step
- ▶ Weights the recent squared gradients more strongly than early ones

Equation 11-7. RMSProp algorithm

1.  $\mathbf{s} \leftarrow \beta \mathbf{s} + (1 - \beta) \nabla_{\theta} J(\theta) \otimes \nabla_{\theta} J(\theta)$
2.  $\theta \leftarrow \theta - \eta \nabla_{\theta} J(\theta) \oslash \sqrt{\mathbf{s} + \epsilon}$

```
optimizer = keras.optimizers.RMSprop(lr=0.001, rho=0.9)
```



# Faster Optimizers: Adam and Nadam

- ▶ Adaptive Moment Estimation
- ▶ Combines Features from Momentum (1) and RMSProp (2 & 5)
- ▶ Steps 3&4 help to boost m & s at start of training

## Equation 11-8. Adam algorithm

1.  $\mathbf{m} \leftarrow \beta_1 \mathbf{m} + (1 - \beta_1) \nabla_{\theta} J(\theta)$
2.  $\mathbf{s} \leftarrow \beta_2 \mathbf{s} + (1 - \beta_2) \nabla_{\theta} J(\theta) \otimes \nabla_{\theta} J(\theta)$
3.  $\widehat{\mathbf{m}} \leftarrow \frac{\mathbf{m}}{1 - \beta_1^t}$
4.  $\widehat{\mathbf{s}} \leftarrow \frac{\mathbf{s}}{1 - \beta_2^t}$
5.  $\theta \leftarrow \theta - \eta \widehat{\mathbf{m}} \oslash \sqrt{\widehat{\mathbf{s}} + \epsilon}$

- ▶ More Optimizers:
  - ▶ Variants: AdaMax and Nadam
  - ▶ Hessian vs Jacobian Methods
  - ▶ [Tensorflow Model Optimization Toolkit](#)

# Learning Rate Scheduling

- ▶ Many options
  - ▶ Power Decay
  - ▶ Exponential
  - ▶ Piecewise Constant
  - ▶ Performance-based
  - ▶ 1cycle (Smith 2018)

```
optimizer = keras.optimizers.SGD(lr=0.01, decay=1e-4)
```

```
lr_scheduler = keras.callbacks.LearningRateScheduler(exponential_decay_fn)
history = model.fit(X_train_scaled, y_train, epochs=n_epochs,
                    validation_data=(X_valid_scaled, y_valid),
                    callbacks=[lr_scheduler])
```

## Why Learning Rate Scheduling?




Figure 11-8. Learning curves for various learning rates  $\eta$

## Module: `tf.keras.optimizers.schedules`

```
reduce_lr = ReduceLRonPlateau(monitor='val_loss', factor=0.2,
                              patience=5, min_lr=0.001)
model.fit(X_train, Y_train, callbacks=[reduce_lr])
```

# Avoid Overfitting Through Regularization

*With four parameters I can fit an elephant and with five I can make him wiggle his trunk.*

—John von Neumann, cited by Enrico Fermi in *Nature* 427

- ▶ DNNs can have millions/billions of parameters
- ▶ How can we prevent overfitting?
- ▶ We've already seen some forms of regularization
  - ▶ Early Stopping
  - ▶ Batch Normalization

# L1 and L2 Regularization

## L2 at Google Devs ML Crash Course

$\text{minimize}(\text{Loss}(\text{Data}|\text{Model}))$

$\text{minimize}(\text{Loss}(\text{Data}|\text{Model}) + \text{complexity}(\text{Model}))$

$L_2 \text{ regularization term} = ||\mathbf{w}||_2^2 = w_1^2 + w_2^2 + \dots + w_n^2$

```
from functools import partial

RegularizedDense = partial(keras.layers.Dense,
                           activation="elu",
                           kernel_initializer="he_normal",
                           kernel_regularizer=keras.regularizers.l2(0.01))

model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    RegularizedDense(300),
    RegularizedDense(100),
    RegularizedDense(10, activation="softmax")
])
model.compile(loss="sparse_categorical_crossentropy", optimizer="nadam", metrics=["accuracy"])
n_epochs = 2
history = model.fit(X_train_scaled, y_train, epochs=n_epochs,
                    validation_data=(X_valid_scaled, y_valid))
```

## Regularizers in tf.keras

Module: tf.keras.regularizers



TensorFlow 1 version

Public API for tf.keras.regularizers namespace.

### Classes

**class L1** : A regularizer that applies a L1 regularization penalty.

**class L1L2** : A regularizer that applies both L1 and L2 regularization penalties.

**class L2** : A regularizer that applies a L2 regularization penalty.

**class Regularizer** : Regularizer base class.

**class l1** : A regularizer that applies a L1 regularization penalty.

**class l2** : A regularizer that applies a L2 regularization penalty.

# Dropout

- ▶ Randomly break edges in the neural network
- ▶ Dropout probability  $p$
- ▶ Analogy: Employee Attendance
- ▶ Implement yes/no Layer by Layer
- ▶ Test time
  - ▶ Activate all
  - ▶ Multiply activations by  $(1-p)$
- ▶ Slower Convergence




Figure 11-9. With dropout regularization, at each training iteration a random subset of all neurons in one or more layers—except the output layer—are “dropped out”; these neurons output 0 at this iteration (represented by the dashed arrows)

```
keras.layers.Dropout(rate=0.2),
```

# Monte Carlo Dropout

```
y_probas = np.stack([model(X_test_scaled, training=True)
                     for sample in range(100)])
y_proba = y_probas.mean(axis=0)
```

- ▶ Gal and Ghahramani, 2016: [Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning](#)
- ▶ training=True
  - ▶ Break connections during inference also!
- ▶ Average
- ▶ Ensemble of weak learners can outperform one strong learner

```
class MCDropout(keras.layers.Dropout):
    def call(self, inputs):
        return super().call(inputs, training=True)
```

# Max Norm Regularization

- ▶ Clipping size of weights, if necessary, after each training step

```
layer = keras.layers.Dense(100, activation="selu", kernel_initializer="lecun_normal",  
.....kernel_constraint=keras.constraints.max_norm(1.))
```

# Summary and Practical Guidelines

- Generally recommended starting hyperparameter configurations:

## Default DNN configuration

Table 11-3. Default DNN configuration

Hyperparameter	Default value
Kernel initializer	He initialization
Activation function	ELU
Normalization	None if shallow; Batch Norm if deep
Regularization	Early stopping (+ $\ell_2$ reg. if needed)
Optimizer	Momentum optimization (or RMSProp or Nadam)
Learning rate schedule	1 cycle

## Self-normalizing NN configuration

Table 11-4. DNN configuration for a self-normalizing net

Hyperparameter	Default value
Kernel initializer	LeCun initialization
Activation function	SELU
Normalization	None (self-normalization)
Regularization	Alpha dropout if needed
Optimizer	Momentum optimization (or RMSProp or Nadam)
Learning rate schedule	1 cycle